



UNIVERSITY
of York

Fundamentals of Neural Networks

Cameron Kyle-Davidson

University of York

This isn't a deep dive into theory!

- Goal: Equip you with enough practical knowledge to build neural networks with PyTorch.
- With a 'good enough' understanding of what's going on behind the scenes.

This Session Structure

- Neural Network Fundamentals
- Practical Session (Getting started with Pytorch)

What actually is a neural network?

Problem Statement

Categorise dogs and cats?




0

1

What actually is a neural network?



↑
Numbers




0

1

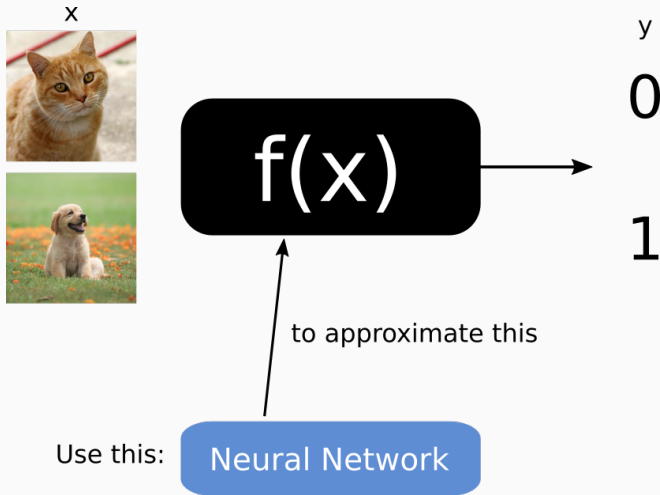
↑
Numbers

What actually is a neural network?




What actually is a neural network?

- Good luck doing that by hand.

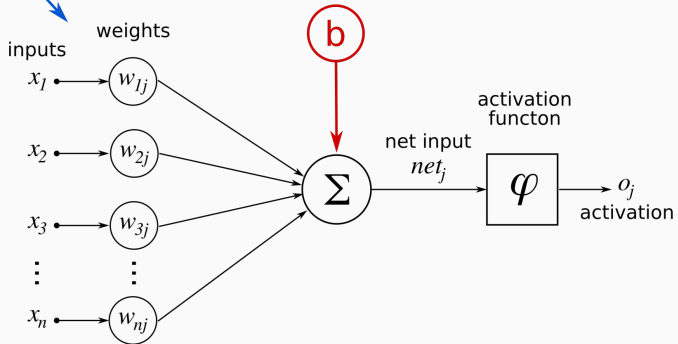



- Luckily, neural networks are **universal function approximators**.

What actually is a neural network?

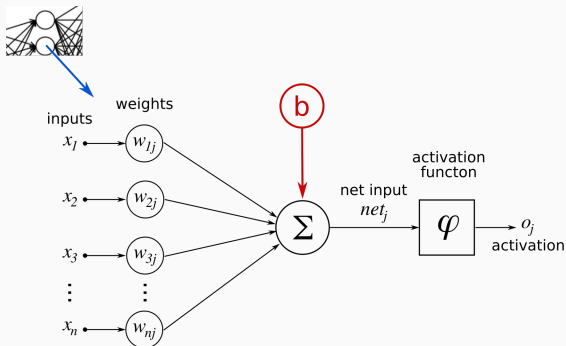


The Artificial Neuron



$$y = \varphi\left(\sum_{k=0}^m w_{kj}x_k + b_j\right)$$

The Artificial Neuron




- **Weights:** A transformation applied to the input to the neuron. Most of what the neural network 'learns' is encoded in the weights.
- **Bias:** Required for the network to describe a variety of functions.

Practical Differences

Don't Worry

Modern ML Libraries tend to abstract away the weights & biases.



```
hidden = nn.Linear(in_features=4, out_features=5)
```

Practical Differences


But they are there.

```
class Linear(Module):
    def __init__(
) -> None:
    factory_kwargs = {"device": device, "dtype": dtype}
    super().__init__()
    self.in_features = in_features
    self.out_features = out_features
    self.weight = Parameter(
        torch.empty((out_features, in_features), **factory_kwargs)
    )
    if bias:
        self.bias = Parameter(torch.empty(out_features, **factory_kwargs))
    else:
        self.register_parameter(name="bias", param=None)
```


* (You're unlikely to see this again, unless you decide to define your own NN layers)

Activation Functions


- Most real-world ‘interesting’ functions aren’t **linear**
- An activation function introduces non-linearity by operating on the output of each neuron.




ReLU



LeakyReLU

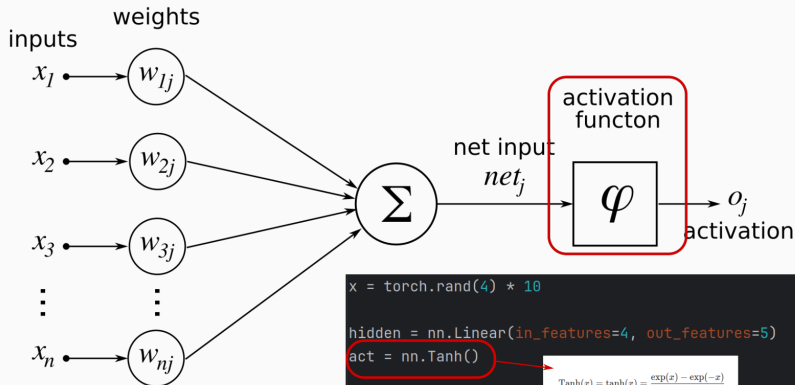


Sigmoid



Tanh

Activation Functions



```
x = torch.rand(4) * 10  
  
hidden = nn.Linear(in_features=4, out_features=5)  
act = nn.Tanh()  
  
pre_act = hidden(x)  
# [ 0.3864, -2.1247,  3.0473,  2.4152, -2.6261]  
post_act = act(pre_act)  
# [ 0.3683, -0.9719,  0.9955,  0.9842, -0.9896]
```

$\text{Tanh}(x) = \tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$

Selection of output-layer activations depends on the task! E.g,
Sigmoid: Binomial Classification, Softmax: Multinomial Classification,
or even linear

Training a Neural Network

Measuring Error


Ground
Truth



$[1, 0, 1]$



Predicted



$[0, 1, 0]$

We need a way to measure error in the NN predictions: a **loss function**.

Loss Functions

There's a lot of loss functions

They tend to be task specific.

`nn.CosineEmbeddingLoss`

`nn.MSELoss`

`nn.L1Loss`

`nn.CrossEntropyLoss`

`nn.NLLLoss`

`nn.BCELoss`

```
y_true = torch.tensor([1.0, 0.0, 1.0])
y_pred = torch.tensor([1.0, 0.0, 1.0])
loss = nn.BCELoss() # 0.0 (perfect)

y_true = torch.tensor([1.0, 0.0, 1.0])
y_pred = torch.tensor([0.0, 1.0, 0.0])
loss = nn.BCELoss() # 100.0 (poor)
```




In practice: have to do some research to find the best loss for your specific problem.

Correcting for Error

The goal of training a NN:

Is to **minimise** the loss function.


- Make the error between predictions and reality as small as possible.
- The **weights** control the loss - so we need to change the weights to reduce the loss.



Backpropagation


We want to go downhill - but which way is that?

- Backpropagation calculates the gradient of the loss function w.r.t the weights.
- Simply: it tells each weight in the network which direction it needs to move (up or down) to reduce the loss.



Optimisers

- Optimisers control *how* to go downhill.
- E.g, How big a step to take



There's many available: SGD, Adam, RMSProp, Adagrad, and more, each with different dynamics.

Practical considerations

Batches:

- Most datasets can't fit into memory
- Efficient parallelisation

Epochs:

- One pass through the entire training set

GPUs?

- Neural networks are essentially matrix multiplication - which GPUs are very good at.
- Massively parallel.

Practical Notebook
